

Docker Persistent Logging & Jupyter Lab Exercises

Exercise 1: Persistent Data Logging

The goal is to run a container that generates log files and ensures those logs are not lost when the container is removed.

1. Launch a Container with a Data Volume:

Start an Alpine Linux container. While launching, mount a directory from your host machine called `log_storage` to the `/var/log` directory inside the container. Provide a screenshot of the docker run command you used.

2. Create a Logging Script:

Inside the running container, create a script at `/usr/local/bin/data_logger.sh`. This script should loop indefinitely, writing a new line with the current date and a "system status OK" message to `/var/log/app.log` every two seconds. Show the contents of your script with the `cat` command and provide a screenshot.

3. Run in Background:

Execute the `data_logger.sh` script in the background. Then, detach from the container's interactive session without stopping it. Capture a screenshot of your terminal showing the commands used to achieve this.

4. Verify on Host:

On your local machine (the host), use the `tail -f` command to monitor the `app.log` file within your `log_storage` directory. This will prove that the data written inside the container is persisting on your host. Provide a screenshot of the `tail` command's output showing the log entries appearing in real-time.

Exercise 2: Running Jupyter Lab in a Container

This task involves running a Jupyter Lab environment in a Docker container and making it accessible from your local machine.

1. Find an Image:

Search Docker Hub for an image such as `jupyter/base-notebook` or `jupyter/scipy-notebook`.

2. Run with Port Forwarding:

Run a container from the Jupyter image. Map port 8888 inside the container to port 9090 on your host machine. The command structure for port mapping is:

```
$> docker run -p 9090:8888 jupyter/base-notebook
```

3. Access Jupyter Lab:

Once the container is running, copy the URL with the token from the terminal output. Then open your web browser and navigate to `http://localhost:9090/?token=` to access Jupyter Lab.

4. Verify Operation:

Create a new notebook, run a simple Python command (like `print("Hello, Docker!")`), and provide a screenshot of your browser showing Jupyter Lab running successfully.

Exercise 3: Configurable Logging via Environment Variables

Here, you'll modify the script from Exercise 1 to make its behavior configurable when the container starts.

1. Parameterize the Script:

Rewrite your `data_logger.sh` script. Instead of a fixed message, it should now read a message prefix from an environment variable called `LOG_PREFIX`. The new log line should look like: `[DATE]: [LOG_PREFIX] - system status OK`.

2. Launch with an Environment Variable:

Run your container, this time setting the `LOG_PREFIX` environment variable to `WebApp1`. The syntax is:

```
$> docker run --env =
```

3. Verify Output:

Check the `app.log` file on your host machine to confirm that the new log lines correctly include the "WebApp1" prefix.

Exercise 4: Building a Robust Script

This task focuses on making your script more reliable by handling missing configuration.

1. Add a Default Value:

Modify `data_logger.sh` again. If the `LOG_PREFIX` environment variable is not set when the container runs, the script should use a default value of `GENERIC`.

2. Test the Fallback:

Prove that your modification works. First, run the container without setting the `LOG_PREFIX` variable and show that the log file contains the "GENERIC" prefix. Then, run it again with an invalid value (e.g., `--env LOG_PREFIX=""`) and show that the default still applies.

Exercise 5: Creating a Custom Docker Image

The final exercise is to package your script into a new, portable Docker image.

1. Create a Dockerfile:

In a new directory, create a file named `Dockerfile`. This file should contain instructions to:

- Start from the alpine base image.
- Copy your final, robust `data_logger.sh` script from your local machine into the image's `/usr/local/bin/` directory.
- Make the script executable.
- Set `data_logger.sh` as the default command to run when a container is started from this image.

2. Build the Image:

Use the `docker build` command to create a new image from your `Dockerfile`. Tag this image as `custom-logger:latest`.

3. Run Your Custom Image:

Launch a container from your new `custom-logger` image. Verify that it automatically starts logging to a mounted volume as intended. Provide a screenshot of your `Dockerfile`'s contents and the `docker run` command used to test your new image.